

housingproject_redux

2024-05-06

```
#housing project redux
#set directory
setwd("/Users/seanmilligan/Desktop/EC424/Homework/housingproject_redux")

#loading packages
pacman::p_load(tidyverse, tidymodels, skimr, caret, leaps, magrittr, janitor, glmnet, zoo)

#loading data
training_df = read.csv('train.csv')
test_df = read.csv('test.csv')

#create age variable (year sold - year built)
house_df = training_df %>% transmute(
  id = Id,
  sale_price = log(SalePrice),
  age = YrSold - YearBuilt,
  remod = YrSold - YearRemodAdd,
  area = GrLivArea,
  lot_area = LotArea,
  cond = OverallCond,
  veneer = MasVnrArea,
  bsmt_sf = TotalBsmtSF,
  bath = FullBath,
  bed_abv = BedroomAbvGr,
  kit_abv = KitchenAbvGr,
  rms_abv = TotRmsAbvGrd,
  fire = Fireplaces,
  grg_age = YrSold - GarageYrBlt,
  wd_dck = WoodDeckSF,
  cl_prch = EnclosedPorch,
  pool = PoolArea
)

#5-fold cross validation
#set seed
set.seed(1234)
#5-fold CV on training dataset
house_cv = house_df %>% vfold_cv(v = 5)
#view CV
house_cv %>% tidy()
```

```
## # A tibble: 7,300 x 3
##   Row Data Fold
```

```
##      <int> <chr>      <chr>
## 1      1 Analysis Fold1
## 2      1 Analysis Fold2
## 3      1 Analysis Fold4
## 4      1 Analysis Fold5
## 5      2 Analysis Fold1
## 6      2 Analysis Fold2
## 7      2 Analysis Fold4
## 8      2 Analysis Fold5
## 9      3 Analysis Fold1
## 10     3 Analysis Fold2
## # i 7,290 more rows
```

```
#define a recipe_all is
recipe_all = recipe(sale_price ~ ., data = house_df)

#putting it together
house_recipe = recipe_all %>%
  #mean imputation for numeric predictors
  step_impute_mean(all_predictors() & all_numeric()) %>%
  #KNN imputation for categorical predictors
  step_impute_knn(all_predictors() & all_nominal(), neighbors = 5 ) %>%
  #create dummies for categorical variables
  step_dummy(all_predictors() & all_nominal())

#putting it together (again for Forward selection)
house_clean = recipe_all %>%
  #mean imputation for numeric predictors
  step_impute_mean(all_predictors() & all_numeric()) %>%
  #KNN imputation for categorical predictors
  step_impute_knn(all_predictors() & all_nominal(), neighbors = 5 ) %>%
  #create dummies for categorical variables
  step_dummy(all_predictors() & all_nominal()) %>%
  #prep and juicing!
  prep() %>% juice()
```

```
#defining model (model type and desired engine)
model_lm =
  linear_reg() %>%
  set_mode('regression') %>%
  set_engine('lm')

#estimating linear regression
#fitting simple linear regression using tidymodels
lm_workflow =
  workflow() %>%
  add_model(model_lm) %>%
  add_recipe(house_recipe)

#fit workflow to data
lm_fit =
  lm_workflow %>%
  fit(data = house_df)
```

```
#view model summary
lm_fit %>% extract_fit_parsnip() %>% tidy()
```

```
## # A tibble: 18 x 5
##   term          estimate std.error statistic  p.value
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)  11.2      0.0435     258.      0
## 2 id          -0.00000482 0.0000108   -0.445 6.57e- 1
## 3 age         -0.00500    0.000298   -16.8  1.02e-57
## 4 remod       -0.00180    0.000335    -5.37  9.25e- 8
## 5 area        0.000272  0.0000196    13.9  2.98e-41
## 6 lot_area    0.00000191 0.000000492   3.88  1.10e- 4
## 7 cond        0.0629    0.00501     12.6  2.28e-34
## 8 veneer      0.0000635  0.0000290    2.19  2.89e- 2
## 9 bsmt_sf     0.000172  0.0000131    13.1  3.35e-37
## 10 bath       0.0339    0.0126      2.69  7.31e- 3
## 11 bed_abv    -0.0299    0.00810    -3.69  2.33e- 4
## 12 kit_abv    -0.144    0.0232     -6.23  6.00e-10
## 13 rms_abv     0.0286    0.00605     4.73  2.48e- 6
## 14 fire       0.0761    0.00852     8.92  1.36e-18
## 15 grg_age    -0.0000751  0.000331    -0.227 8.21e- 1
## 16 wd_dck     0.000118  0.0000390     3.02  2.58e- 3
## 17 cl_prch    0.000285  0.0000821     3.47  5.40e- 4
## 18 pool      -0.000474  0.000117    -4.06  5.12e- 5
```

```
#fitting linear regression w/ 5-fold CV
```

```
fit_lm_cv =
  workflow() %>%
  add_model(model_lm) %>%
  add_recipe(house_recipe) %>%
  fit_resamples(house_cv)
#checking performance
fit_lm_cv %>% collect_metrics()
```

```
## # A tibble: 2 x 6
##   .metric .estimator mean      n std_err .config
##   <chr>   <chr>    <dbl> <int>  <dbl> <chr>
## 1 rmse    standard  0.179     5  0.0164 Preprocessor1_Model11
## 2 rsq     standard  0.799     5  0.0386 Preprocessor1_Model11
```

```
#checking performance within each fold
```

```
fit_lm_cv %>% collect_metrics(summarize = F)
```

```
## # A tibble: 10 x 5
##   id      .metric .estimator .estimate .config
##   <chr> <chr>    <chr>      <dbl> <chr>
## 1 Fold1 rmse    standard    0.159 Preprocessor1_Model11
## 2 Fold1 rsq     standard    0.844 Preprocessor1_Model11
## 3 Fold2 rmse    standard    0.147 Preprocessor1_Model11
## 4 Fold2 rsq     standard    0.856 Preprocessor1_Model11
## 5 Fold3 rmse    standard    0.200 Preprocessor1_Model11
## 6 Fold3 rsq     standard    0.771 Preprocessor1_Model11
```

```
## 7 Fold4 rmse      standard      0.234 Preprocessor1_Model1
## 8 Fold4 rsq       standard      0.659 Preprocessor1_Model1
## 9 Fold5 rmse      standard      0.156 Preprocessor1_Model1
## 10 Fold5 rsq      standard      0.863 Preprocessor1_Model1
```

```
#forward selection for available variables
```

```
train_forward1 = train(
  y = house_clean[["sale_price"]],
  x = house_clean %>% dplyr::select(-sale_price),
  trControl = trainControl(method = "cv", number = 5),
  method = "leapForward",
  tuneGrid = expand.grid(nvmax = 1:18)
)
```

```
## Warning: Setting row names on a tibble is deprecated.
## Setting row names on a tibble is deprecated.
## Setting row names on a tibble is deprecated.
## Setting row names on a tibble is deprecated.
## Setting row names on a tibble is deprecated.
## Setting row names on a tibble is deprecated.
```

```
train_forward1$results
```

##	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	0.2842110	0.4940369	0.2090159	0.02306191	0.02998926	0.01160796
## 2	2	0.2175185	0.7053151	0.1556852	0.03166460	0.05487518	0.01582137
## 3	3	0.2077624	0.7322241	0.1418299	0.04111959	0.07660834	0.01195041
## 4	4	0.1875634	0.7813858	0.1233449	0.04395227	0.08041279	0.01282676
## 5	5	0.1814070	0.7947432	0.1202967	0.04273816	0.07580137	0.01250764
## 6	6	0.1783330	0.8017521	0.1193154	0.04223832	0.07343004	0.01236017
## 7	7	0.1810199	0.7962162	0.1192956	0.03965631	0.06787574	0.01000858
## 8	8	0.1804600	0.7973894	0.1182092	0.04066592	0.06951185	0.01042072
## 9	9	0.1806877	0.7970194	0.1178035	0.04146419	0.07121055	0.01033003
## 10	10	0.1802968	0.7979378	0.1179301	0.04017025	0.06864131	0.01011049
## 11	11	0.1808203	0.7971098	0.1180978	0.04017259	0.06849221	0.01013907
## 12	12	0.1795220	0.7998376	0.1175321	0.04030441	0.06890698	0.01044305
## 13	13	0.1785452	0.8017009	0.1172972	0.04046780	0.06900410	0.01105157
## 14	14	0.1778665	0.8031226	0.1166371	0.04134069	0.07068748	0.01136296
## 15	15	0.1769187	0.8053560	0.1156569	0.04136260	0.07068315	0.01114497
## 16	16	0.1773657	0.8043758	0.1158068	0.04174652	0.07159018	0.01105627
## 17	17	0.1773501	0.8044076	0.1157926	0.04174004	0.07158181	0.01107723
## 18	18	0.1773501	0.8044076	0.1157926	0.04174004	0.07158181	0.01107723

```
#model with all variables has lowest RMSE ^^^
```

Our model containing all variables possesses the lowest Residual Mean Squared Error.

```
#using lasso regression w/tidy models with CV
```

```
#standardizing data for use
```

```
house_recipe_lasso = house_clean %>% recipe(sale_price ~ .) %>%
  update_role(id, new_role = 'id_variable') %>%
  step_normalize(all_predictors() & all_numeric()) %>%
```

```

step_dummy(all_predictors() & all_nominal()) %>%
step_rename_at(everything(), fn = str_to_lower)
#time to juice it up
house_recipe_lasso_clean = house_recipe %>% prep() %>% juice()

#using lasso and ridge w/5-fold cross validation for penalty on lasso and regression
set.seed(12345)
ctrl_cv = trainControl(method = "cv", number = 5)

#define range of lambdas (glmnet wants decreasing range)
lambdas = 10^seq(from = 5, to = -2, length = 100)

#defining model
lasso_est = linear_reg(penalty = tune(), mixture = 1) %>% set_engine('glmnet')

#defining lasso workflow
workflow_lasso = workflow() %>%
  add_model(lasso_est) %>% add_recipe(house_recipe_lasso)
#CV w/range of lambdas
cv_lasso =
  workflow_lasso %>%
  tune_grid(
    resamples = vfold_cv(house_clean, v = 5),
    grid = data.frame(penalty = lambdas),
    metrics = metric_set(rmse)
  )
#show best models
cv_lasso %>% show_best()

```

```
## Warning in show_best(.): No value of 'metric' was given; "rmse" will be used.
```

```
## # A tibble: 5 x 7
##   penalty .metric .estimator mean      n std_err .config
##   <dbl> <chr>   <chr>   <dbl> <int>   <dbl> <chr>
## 1  0.01  rmse     standard 0.181     5  0.0146 Preprocessor1_Model001
## 2  0.0118 rmse     standard 0.181     5  0.0143 Preprocessor1_Model002
## 3  0.0138 rmse     standard 0.182     5  0.0139 Preprocessor1_Model003
## 4  0.0163 rmse     standard 0.183     5  0.0134 Preprocessor1_Model004
## 5  0.0192 rmse     standard 0.185     5  0.0128 Preprocessor1_Model005
```

```

#finding best lambda
cv_lasso$.metrics

```

```

## [[1]]
## # A tibble: 100 x 5
##   penalty .metric .estimator .estimate .config
##   <dbl> <chr>   <chr>         <dbl> <chr>
## 1  0.01  rmse     standard      0.238 Preprocessor1_Model001
## 2  0.0118 rmse     standard      0.237 Preprocessor1_Model002
## 3  0.0138 rmse     standard      0.237 Preprocessor1_Model003
## 4  0.0163 rmse     standard      0.236 Preprocessor1_Model004
## 5  0.0192 rmse     standard      0.235 Preprocessor1_Model005

```

```

## 6 0.0226 rmse standard 0.234 Preprocessor1_Model006
## 7 0.0266 rmse standard 0.233 Preprocessor1_Model007
## 8 0.0313 rmse standard 0.233 Preprocessor1_Model008
## 9 0.0368 rmse standard 0.233 Preprocessor1_Model009
## 10 0.0433 rmse standard 0.234 Preprocessor1_Model010
## # i 90 more rows
##
## [[2]]
## # A tibble: 100 x 5
##   penalty .metric .estimator .estimate .config
##   <dbl> <chr> <chr> <dbl> <chr>
## 1 0.01 rmse standard 0.159 Preprocessor1_Model001
## 2 0.0118 rmse standard 0.159 Preprocessor1_Model002
## 3 0.0138 rmse standard 0.161 Preprocessor1_Model003
## 4 0.0163 rmse standard 0.162 Preprocessor1_Model004
## 5 0.0192 rmse standard 0.164 Preprocessor1_Model005
## 6 0.0226 rmse standard 0.166 Preprocessor1_Model006
## 7 0.0266 rmse standard 0.170 Preprocessor1_Model007
## 8 0.0313 rmse standard 0.174 Preprocessor1_Model008
## 9 0.0368 rmse standard 0.180 Preprocessor1_Model009
## 10 0.0433 rmse standard 0.186 Preprocessor1_Model010
## # i 90 more rows
##
## [[3]]
## # A tibble: 100 x 5
##   penalty .metric .estimator .estimate .config
##   <dbl> <chr> <chr> <dbl> <chr>
## 1 0.01 rmse standard 0.174 Preprocessor1_Model001
## 2 0.0118 rmse standard 0.175 Preprocessor1_Model002
## 3 0.0138 rmse standard 0.176 Preprocessor1_Model003
## 4 0.0163 rmse standard 0.177 Preprocessor1_Model004
## 5 0.0192 rmse standard 0.179 Preprocessor1_Model005
## 6 0.0226 rmse standard 0.181 Preprocessor1_Model006
## 7 0.0266 rmse standard 0.184 Preprocessor1_Model007
## 8 0.0313 rmse standard 0.188 Preprocessor1_Model008
## 9 0.0368 rmse standard 0.193 Preprocessor1_Model009
## 10 0.0433 rmse standard 0.197 Preprocessor1_Model010
## # i 90 more rows
##
## [[4]]
## # A tibble: 100 x 5
##   penalty .metric .estimator .estimate .config
##   <dbl> <chr> <chr> <dbl> <chr>
## 1 0.01 rmse standard 0.160 Preprocessor1_Model001
## 2 0.0118 rmse standard 0.162 Preprocessor1_Model002
## 3 0.0138 rmse standard 0.163 Preprocessor1_Model003
## 4 0.0163 rmse standard 0.165 Preprocessor1_Model004
## 5 0.0192 rmse standard 0.168 Preprocessor1_Model005
## 6 0.0226 rmse standard 0.171 Preprocessor1_Model006
## 7 0.0266 rmse standard 0.175 Preprocessor1_Model007
## 8 0.0313 rmse standard 0.181 Preprocessor1_Model008
## 9 0.0368 rmse standard 0.188 Preprocessor1_Model009
## 10 0.0433 rmse standard 0.193 Preprocessor1_Model010
## # i 90 more rows

```

```
##
## [[5]]
## # A tibble: 100 x 5
##   penalty .metric .estimator .estimate .config
##   <dbl> <chr>   <chr>         <dbl> <chr>
## 1 0.01   rmse     standard     0.173 Preprocessor1_Model001
## 2 0.0118 rmse     standard     0.174 Preprocessor1_Model002
## 3 0.0138 rmse     standard     0.176 Preprocessor1_Model003
## 4 0.0163 rmse     standard     0.177 Preprocessor1_Model004
## 5 0.0192 rmse     standard     0.180 Preprocessor1_Model005
## 6 0.0226 rmse     standard     0.182 Preprocessor1_Model006
## 7 0.0266 rmse     standard     0.186 Preprocessor1_Model007
## 8 0.0313 rmse     standard     0.191 Preprocessor1_Model008
## 9 0.0368 rmse     standard     0.196 Preprocessor1_Model009
## 10 0.0433 rmse     standard     0.201 Preprocessor1_Model010
## # i 90 more rows
```

```
#lowest RMSE ~0.182 @ lambda = 0.0118
```

```
#fitting final model
```

```
final_lasso = glmnet(
  x = house_clean %>% dplyr::select(-sale_price, -id) %>% as.matrix(),
  y = house_clean$sale_price,
  standardize = T,
  alpha = 1,
  lambda = 0.0118
)
```

```
#cleaning test data set
```

```
#create age variable (year sold - year built)
```

```
pred_df = test_df %>% transmute(
  id = Id,
  age = YrSold - YearBuilt,
  remod = YrSold - YearRemodAdd,
  area = GrLivArea,
  lot_area = LotArea,
  cond = OverallCond,
  veneer = MasVnrArea,
  bsmt_sf = TotalBsmtSF,
  bath = FullBath,
  bed_abv = BedroomAbvGr,
  kit_abv = KitchenAbvGr,
  rms_abv = TotRmsAbvGrd,
  fire = Fireplaces,
  grg_age = YrSold - GarageYrBlt,
  wd_dck = WoodDeckSF,
  cl_prch = EnclosedPorch,
  pool = PoolArea
)
```

```
#create function to remove NA values from all columns
```

```
rep_NA_func = function(data) {
  for (col in names(data)) {
    data[[col]] = na.aggregate(data[[col]])
  }
}
```

```

    return(data)
}

#clean prediction dataframe w/rep_NA function
pred_clean = rep_NA_func(pred_df)

#logarithmic prediction
pred_log = predict(
  final_lasso,
  type = "response",
  #our chosen lambda
  s = 0.0118,
  #our data
  newx = pred_clean %>% dplyr::select(-id) %>% as.matrix()
)

#final prediction w/o logarithms
pred_final = exp(pred_log)

#create submission dataset
submit_df = data.frame(
  Id = test_df$Id,
  SalePrice = pred_final
)

#change name of s1 to SalePrice
colnames(submit_df)[colnames(submit_df) == 's1'] = 'SalePrice'

#view first few lines of dataset
head(submit_df)

##      Id SalePrice
## 1 1461  114853.9
## 2 1462  142464.9
## 3 1463  195685.8
## 4 1464  206797.7
## 5 1465   162418.2
## 6 1466   186738.9

#save dataset as CSV
write_csv(x = submit_df, file = 'spm_submit_redux.csv')

```